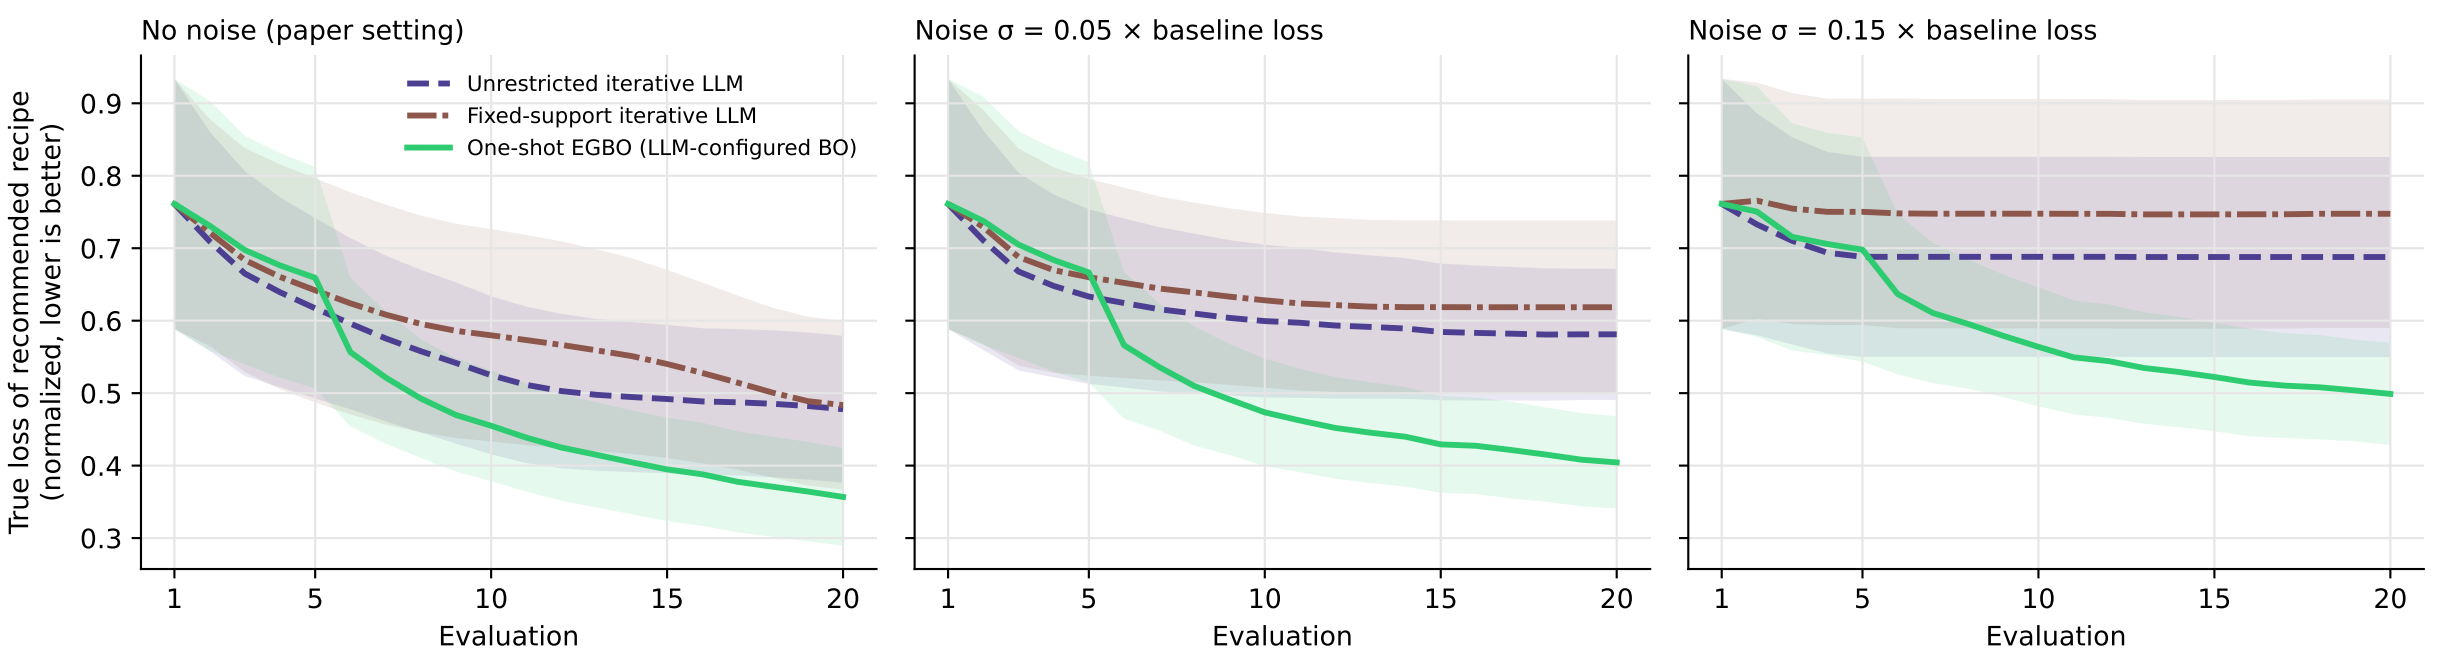
No noise (paper setting)
Noise σ = 0.05 × baseline loss
Noise σ = 0.15 × baseline loss
Unrestricted iterative LLM
Fixed-support iterative LLM
One-shot EGBO (LLM-configured BO)
True loss of recommended recipe
(normalized, lower is better)
0.9
0.8
0.7
0.6
0.5
0.4
0.3
1
5
10
15
20
Evaluation